

50 points total. 70+% correctness (35+ points) is needed to pass. Remember: you must pass all assignments to pass the class. The assignment is due at the beginning of the next class.

1. Profiling and Vectorization (15 points)

In this question we are going to use Python profiling tools to determine the efficiency of using NumPy's vectorization capabilities. To do so we will compute the first 1000 powers (inclusive and starting from 0) of 2 in the following ways. Report the amount of time it takes to run each of these methods 100,000 times.

To profile each of the statements you will create, use the `timeit` module. There are many ways to use this module (see <http://docs.python.org/2.7/library/timeit.html>) but you should only need to use the `timeit()` function. To time multi-line statements, create the statement using triple quotes. Remember to import the necessary modules and create any necessary objects in the setup phase.

- (a) append to an initially empty list sequentially

```
timeit.timeit(stmt="""
a = []
for i in xrange(1000):
    a.append(2 ** i)""", number=100000)
```

366 seconds, very slow

- (b) create an list of zeros of size 1000 and assign elements sequentially

```
timeit.timeit(stmt="""
for i in xrange(len(a)):
    a[i] = 2. ** i""",
    setup = """a = [0] * 1000""", number=100000)
```

356 seconds, pre-allocating the array doesn't help much

- (c) use a list comprehension

```
timeit.timeit(stmt="""[2. ** i for i in xrange(1000)]""", number=100000)
```

355 seconds, list comprehension also does not help much

- (d) use the python map function

```
timeit.timeit(stmt="""map(f, a)""",
    setup="""a=range(1000)
f = lambda i: 2. ** i
""", number=100000)
```

378 seconds, Python's built in vectorization does not help.

- (e) use vectorization (NumPy)

```
timeit.timeit(stmt="""2. ** np.arange(1000)""",
    setup="""import numpy as np""", number=100000)
```

17 seconds, order of magnitude speedup.

2. Copies and Views (10 points)

NumPy arrays are objects and follow the same assignment and copy rules as ordinary python objects. However, when we slice an array python returns a *view* on that same data, meaning that a new object is created but shares the same underlying data. Integer indexing and boolean indexing do not create views. Determine the value of the following statements and *very* briefly explain why.

```
(a) >>> def f(arr):
...     arr[0, 0] = 42
...     return arr
>>> A = np.ones((2, 5))
>>> B = f(A)
>>> A is B
```

True, NumPy arrays are objects and python passes objects by reference. When we call `f(A)` a reference to `A` is passed and returned where we then assign `B` to that same reference. Hence `A` is `B` and the first element of both `A` and `B` is 42.

```
(b) >>> A = np.ones((2, 5))
>>> B = A[:, 1:3]
>>> B[0,0] = 42.
>>> A[0, 1] == 1.
```

False, using slicing operations returns a view, so `B` is a view onto the same underlying data as `A`. When we change the data of `B` we are changing the data of `A`. So `A[0,1]` is now 42. NumPy does this to avoid copying data when possible, but the user needs to be careful of these cases.

```
(c) >>> A = np.ones((2, 5))
>>> B = A[:, np.array([False, True, True, False, False])]
>>> B[0, 0] = 42
>>> A[0, 1] == 1
```

True, only slicing returns views. `B` is not a view onto `A` so changing the data of `B` does not change the data of `A`.

```
(d) >>> A = np.ones((2, 5))
>>> B = A[:, np.array([1, 3])]
>>> B[0, 0] = 42
>>> A[0, 1] == 1
```

True, only slicing returns views.

```
(e) >>> A = np.ones((2, 5))
>>> B = A[:, 1:3]
>>> A += 1
>>> np.all(B == 2)
```

True, `B` is a view onto `A` and `+=` edits `A` in place. If we had instead called `A = A + 1`, then we would be creating a new set of data, and `B` would remain 1's so that the final statement would be false.

3. Kernel Density Estimate (KDE) (10 points)

Given a sample x_1, x_2, \dots, x_n from an unknown distribution f the kernel density estimate of f at point x with kernel K_b is defined as

$$\hat{f}(x; b) = \frac{1}{n} \sum_{i=1}^n K_b(x - x_i)$$

where the kernel must satisfy $\int_{-\infty}^{\infty} K(u) du = 1$ and $K(-u) = K(u)$. The parameter b is known as the bandwidth and controls the width of the kernel used. There are many possible choices for kernels, and we will use the triangular kernel

$$K_b(z) = \frac{1}{b} (1 - |\frac{z}{b}|) * \mathbb{I}(|\frac{z}{b}| \leq 1)$$

Write a two line python function `kde(x, data, bw)` without list comprehension that takes the three parameters below and returns $\hat{f}(x)$, the KDE evaluated at x .

- (a) `x` the point to evaluate the KDE
- (b) `data` the sample of points from f ; above denoted x_1, x_2, \dots, x_n
- (c) `bw` the bandwidth of the kernel

```
kde = lambda x, data, bw: np.mean(1./bw * np.maximum(1.- np.abs(x-data) / bw, 0))
```

Or without using max

```
kde = lambda x, data, bw: np.mean(1./bw * (1.- np.abs(x-data)/bw) * (np.abs(x-data)/bw<=1))
```

4. SciPy Optimization (15 points)

The `scipy.optimize` module contains functions that perform numerical optimization. Use `scipy.optimize.minimize` to minimize the following function.

$$f(x) = \frac{1}{2} x^T A x - b^T x$$

When $A = \begin{pmatrix} 1 & 1 \\ 1 & 2 \end{pmatrix}$ and $b = \begin{pmatrix} 1 \\ 1 \end{pmatrix}$. Verify that your solution x^* is a solution to the system of linear equations defined by A and b . That is, make sure $Ax^* = b$.

```
import scipy.optimize
A = np.array([[1., 1.], [1., 2.]])
b = np.ones(2)
x_star = scipy.optimize.fmin(func=lambda x: .5 * np.dot(x.T, np.dot(A, x)) - np.dot(b, x),
                             x0=np.zeros(2))
```

```
np.dot(A, x_star)
```

$$x^* = \begin{pmatrix} 1 \\ 0 \end{pmatrix}$$

5. matplotlib (0 points)

It will be useful to have matplotlib installed for lecture 5. Please install `matplotlib` on your system and make sure the following lines produce a graph

```
>>> import matplotlib.pyplot as plt
>>> plt.plot(np.arange(10))
>>> plt.show()
```